

Capstone – Spark Project

Erwan Fouillen



1. Introduction

The goal of this second task on spark is to showcase the simultaneous application of several big data technologies, e.g.:

- Data loading
- In-memory distributed computing
- Parallelized SQL (in our case, with Hive and/or Cassandra)
- Containerized applications (use of Docker to run Spark and Cassandra together)

Please note that, as I had issues with AWS credits, I ended up doing everything locally with Docker:

1. Use of a bash script to download data files concurrently (see docker image in resources)
2. Data loading with Spark DataFrames
3. Query the data with SQL statements (Hive) or CQL statements (Cassandra)
4. Store the results in a Cassandra database (in my case, PrestoDB)
5. GitHub repository: <https://github.com/ErwFoui/CloudComputingProject>

2. Data Extraction

As I didn't use AWS to complete this assignment, I directly downloaded the data from the website of the US Bureau of Transportation Statistics. Plainly put, and as shown in the **get_data.sh** file of my github repository, I parallelized in a bash script the download of monthly flight CSVs from 1987 to 2008 as follows:

1. **cURL** https://www.transtats.bts.gov/DL_SelectFields.asp?Table_ID=236 with a SQL statement specifying the required fields plus the **year** and **month** filters. The required fields include those useful to answer the assignment's questions, i.e.: YEAR, MONTH, DAY_OF_MONTH, DAY_OF_WEEK, FL_DATE, OP_UNIQUE_CARRIER, OP_CARRIER_FL_NUM, ORIGIN_AIRPORT_ID, ORIGIN, DEST_AIRPORT_ID, DEST, CRS_DEP_TIME, DEP_TIME, DEP_DELAY, DEP_DELAY_NEW, CRS_ARR_TIME, ARR_TIME, ARR_DELAY, ARR_DELAY_NEW, CANCELLED
2. The server answers by specifying the location of the zipped file for the given year/month filters. **wget** is used to download the file at the URL sent by the server.
3. Zipped files are unzipped and renamed.
4. The 21 * 12 csv files are downloaded in parallel by forking bash processes (see **get_data.sh**):

```
max_num_processes=$(ulimit -u)
limiting_factor=4
num_processes=$((max_num_processes/limiting_factor))

# Download all CSVs in parallel
for year in `seq 1987 2008`
do
  for month in `seq 1 12`
  do
    ((i=i%num_processes)); ((i+==0)) && wait
    download_csv $year $month &
  done
done
```

5. Once download is complete, data is extracted from the CSV files with Spark's DataFrame API, in a SparkContext or a SQLContext.
6. As running Kafka locally proved too slow when sending data to/consuming data from the Kafka, I ended up using directly Spark to extract and concatenate the CSV files in lazily evaluated DataFrames.

3. Architecture

All in all, the architecture is very similar to that offered by AWS, i.e., use of PySpark to distribute files across Resilient Distributed Datasets, then use Hive-like statements to query these files with simple SQL queries, and finally store tables in a database (Cassandra). Everything was done locally

with the help of the following Docker image: <https://github.com/Yannael/kafka-sparkstreaming-cassandra>.

All a user has to do is to start this docker image with a volume linking the host machine's /Data folder to the container's directories, run **get_data.sh** to concurrently download all the files, and finally run any of the 3 Jupyter notebooks to extract and process the data. Please note that, as I was only using my personal laptop, RAM capabilities prevented me from keeping all the CSV files in memory. Consequently, I ended up using datasets from 1995 (instead of 1987) to 2008, hence some discrepancies between results from part 1 and those from part 2.

4. Results

This last part shows answers to the assignment's questions. All queries and results are shown in the **SparkQueries – Q1/2/3** files. Please note that, as there is no missing data in the csv files downloaded at https://www.transtats.bts.gov/DL_SelectFields.asp?Table_ID=236 (contrary to some parts of the AWS snapshot), actual results might slightly change. However, the relative rankings match those indicated on the Capstone project's webpage.

1.1

```
+-----+-----+
|ORIGIN|all_flights|
+-----+-----+
|   ORD   | 9252980 |
|   ATL   | 8885867 |
|   DFW   | 7926248 |
|   LAX   | 5844795 |
|   PHX   | 5047005 |
|   DEN   | 4508601 |
|   IAH   | 4423809 |
|   DTW   | 4197592 |
|   LAS   | 4024215 |
|   MSP   | 3939329 |
+-----+-----+
```

1.2

```
+-----+-----+
|airline|avg_arr_delay|
+-----+-----+
|   HA   | -0.775 |
|   KH   |  1.157 |
|   F9   |  5.693 |
|   WN   |  5.84  |
|   OO   |  5.877 |
|   9E   |  6.108 |
|   TZ   |  6.129 |
|   NW   |  6.298 |
|   US   |  6.471 |
|   DH   |  6.798 |
+-----+-----+
```

1.3

+---+-----+	
dow avg_arr_delay	
+---+-----+	
6	4.281
2	6.008
3	7.172
7	7.249
1	7.297
4	9.341
5	10.257
+---+-----+	

2.1

+---+-----+-----+-----+				
origin	carrier	avg_delay	rank	
+---+-----+-----+-----+				
	CMH	FL	1.948	1
	CMH	DH	3.491	2
	CMH	AA	3.53	3
	CMH	DL	4.535	4
	CMH	NW	4.969	5
	CMH	US	5.66	6
	CMH	YV	7.961	7
	CMH	TW	8.03	8
	CMH	9E	8.326	9
	CMH	WN	8.37	10
	SRQ	TZ	-0.382	1
	SRQ	XE	1.49	2
	SRQ	YV	3.426	3
	SRQ	US	3.489	4
	SRQ	DL	5.31	5
	SRQ	MQ	5.351	6
	SRQ	TW	5.967	7
	SRQ	FL	6.061	8
	SRQ	NW	6.443	9
	SRQ	CO	9.286	10
	BOS	TZ	3.064	1
	BOS	NW	6.951	2
	BOS	DL	7.162	3
	BOS	EV	7.208	4
	BOS	US	7.941	5
	BOS	XE	8.987	6
	BOS	AA	9.178	7
	BOS	UA	9.577	8
	BOS	B6	10.046	9
	BOS	FL	10.131	10
	SEA	OO	2.877	1
	SEA	YV	5.122	2
	SEA	US	5.852	3
	SEA	DL	6.188	4

	SEA	TZ	6.345	5
	SEA	NW	7.009	6
	SEA	CO	7.031	7
	SEA	HA	7.133	8
	SEA	EV	7.896	9
	SEA	AA	7.951	10
	JFK	UA	5.63	1
	JFK	CO	7.656	2
	JFK	XE	8.114	3
	JFK	DH	8.743	4
	JFK	NW	11.032	5
	JFK	B6	11.229	6
	JFK	TW	11.496	7
	JFK	DL	11.728	8
	JFK	US	11.924	9
	JFK	AA	11.939	10
+-----+-----+-----+-----+				

2.2

+-----+-----+-----+-----+				
origin destination avg_delay rank				
+-----+-----+-----+-----+				
	CMH	AUS	-5.0	1
	CMH	OMA	-5.0	2
	CMH	SYR	-5.0	3
	CMH	CLE	0.526	4
	CMH	MSN	1.0	5
	CMH	SLC	3.555	6
	CMH	CLT	3.667	7
	CMH	IAD	4.158	8
	CMH	MEM	4.295	9
	CMH	IAH	4.381	10
	SRQ	IAH	-0.688	1
	SRQ	TPA	-0.56	2
	SRQ	EYW	0.0	3
	SRQ	DFW	1.858	4
	SRQ	FLL	2.0	5
	SRQ	MCO	2.066	6
	SRQ	RSW	2.199	7
	SRQ	CLE	2.462	8
	SRQ	MDW	2.638	9
	SRQ	CLT	2.915	10
	BOS	SWF	-5.0	1
	BOS	ONT	-3.0	2
	BOS	GGG	1.0	3
	BOS	AUS	1.324	4

	BOS	LGA	2.749	5
	BOS	MSY	3.763	6
	BOS	LGB	5.431	7
	BOS	DCA	5.538	8
	BOS	MKE	5.945	9
	BOS	MDW	6.003	10
	SEA	EUG	0.0	1
	SEA	PIH	1.0	2
	SEA	PSC	2.651	3
	SEA	CVG	3.571	4
	SEA	MEM	4.0	5
	SEA	CLE	4.749	6
	SEA	PIT	5.13	7
	SEA	LIH	5.381	8
	SEA	SNA	5.649	9
	SEA	CLT	5.68	10
	JFK	SWF	-10.5	1
	JFK	STX	-2.0	2
	JFK	ABQ	0.0	3
	JFK	ANC	0.0	4
	JFK	ISP	0.0	5
	JFK	MYR	0.0	6
	JFK	BGR	3.86	7
	JFK	BQN	3.95	8
	JFK	CHS	4.403	9
	JFK	AGS	4.75	10
+-----+-----+-----+-----+				

2.3

+-----+-----+-----+-----+-----+				
origin	destination	avg_delay	carrier	rank
+-----+-----+-----+-----+-----+				
	BOS	LGA	0.663	US 1
	BOS	LGA	2.16	DL 2
	BOS	LGA	12.482	MQ 3
	BOS	LGA	25.6	AA 4
	BOS	LGA	30.448	OH 5
	BOS	LGA	133.0	TZ 6
	LGA	BOS	-3.168	US 1
	LGA	BOS	1.167	DL 2
	LGA	BOS	9.471	MQ 3
	LGA	BOS	27.985	OH 4
	LGA	BOS	40.0	AA 5
	MSP	ATL	5.378	OO 1
	MSP	ATL	6.021	DL 2


	MSP		ATL		6.398		FL		3	
	MSP		ATL		7.74		NW		4	
	MSP		ATL		8.352		OH		5	
	MSP		ATL		10.287		EV		6	
	OKC		DFW		1.359		EV		1	
	OKC		DFW		4.049		AA		2	
	OKC		DFW		4.71		MQ		3	
	OKC		DFW		12.835		OO		4	
	OKC		DFW		13.401		DL		5	
	OKC		DFW		47.5		OH		6	
+-----+-----+-----+-----+-----+										



2.4

+-----+-----+-----+		
origin	destination	avg_delay
+-----+-----+-----+		
	BOS	LGA
	LGA	BOS
	MSP	ATL
	OKC	DFW
+-----+-----+-----+		

3.1

Let us compare the three Quantile-Quantile plots below (drawn in python with the **scipy.stats** package). It appears that the distribution that is closest to a normal distribution isn't that of "Log(Log(X))" but "Log(X)", which allows us to conclude that X doesn't follow Zipf 's law.





3.2 In order, the columns are X, Y, Z, first leg's flight date, first leg's airline, first leg's flight number, and then the same pattern applies to the second leg of the trip:

```
+---+---+---+-----+-----+-----+-----+-----+
-+-----+
|  x|  y|  z|xy_fl_date|carrier_xy|flight_xy|yz_fl_date|carrier_yz|flight_y
z|total_delay|
+---+---+---+-----+-----+-----+-----+-----+
-+-----+
|BOS|ATL|LAX|    4/3/08|    FL|    270|    4/5/08|    FL|    4
0|          5|
|PHX|JFK|MSP|    9/7/08|    B6|    178|    9/9/08|    NW|    60
9|          -42|
|DFW|STL|ORD|    1/14/08|    AA|   1336|    1/16/08|    AA|   224
5|          -19|
|LAX|MIA|LAX|    5/16/08|    AA|    280|    5/18/08|    AA|    45
6|          -9|
+---+---+---+-----+-----+-----+-----+-----+
-+-----+
```

Resources:

- Docker image: <https://hub.docker.com/r/yannael/kafka-sparkstreaming-cassandra>
- Zipf's law: https://en.wikipedia.org/wiki/Zipf%27s_law
- Q-Q plots: <https://en.wikipedia.org/wiki/Q%E2%80%93plot>
- GitHub Repository: <https://github.com/ErwFoui/CloudComputingProject>